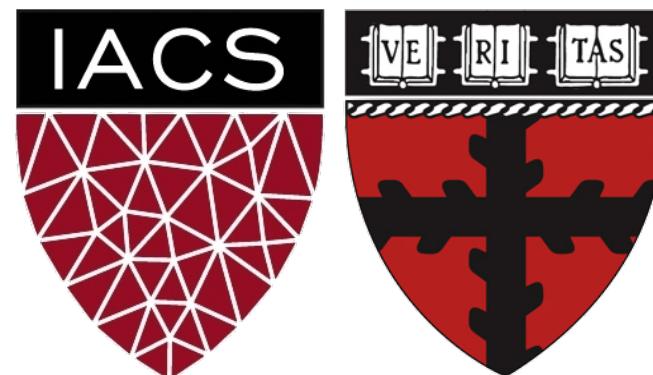


# Transfer Learning

Pavlos Protopapas

Institute for Applied Computational Science  
Harvard





Marios Mattheakis



Robbert Struyven

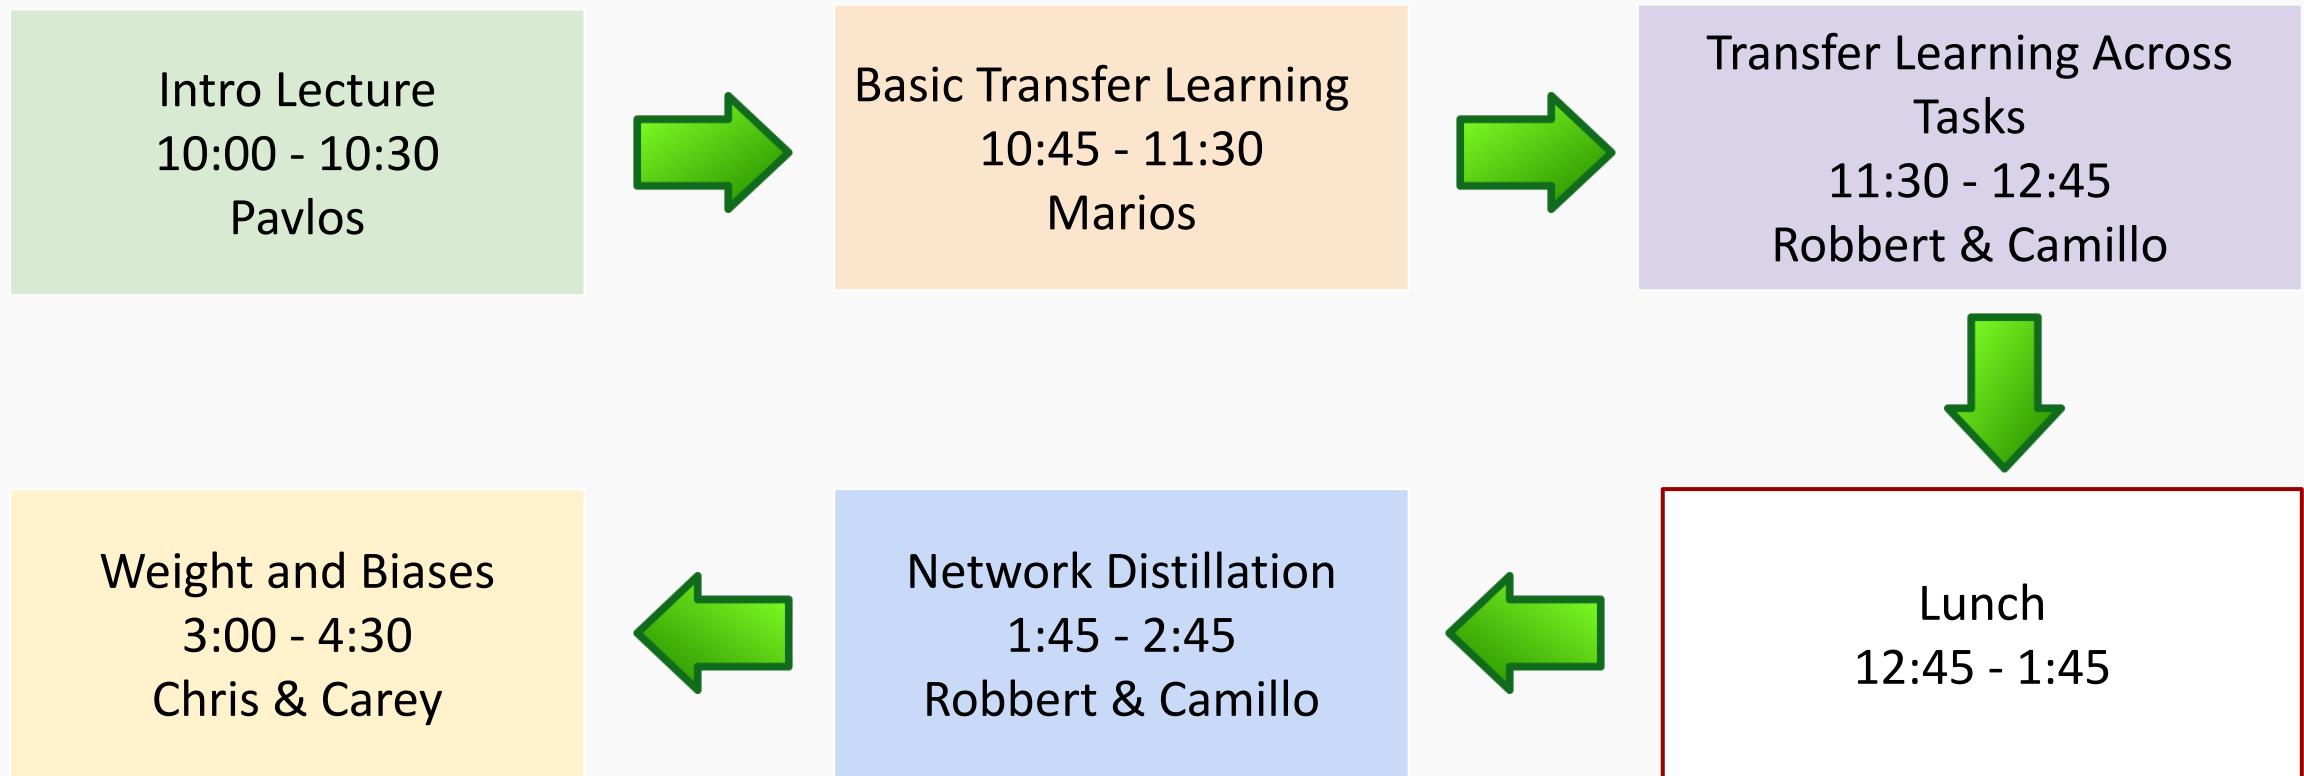


Camilo Fosco

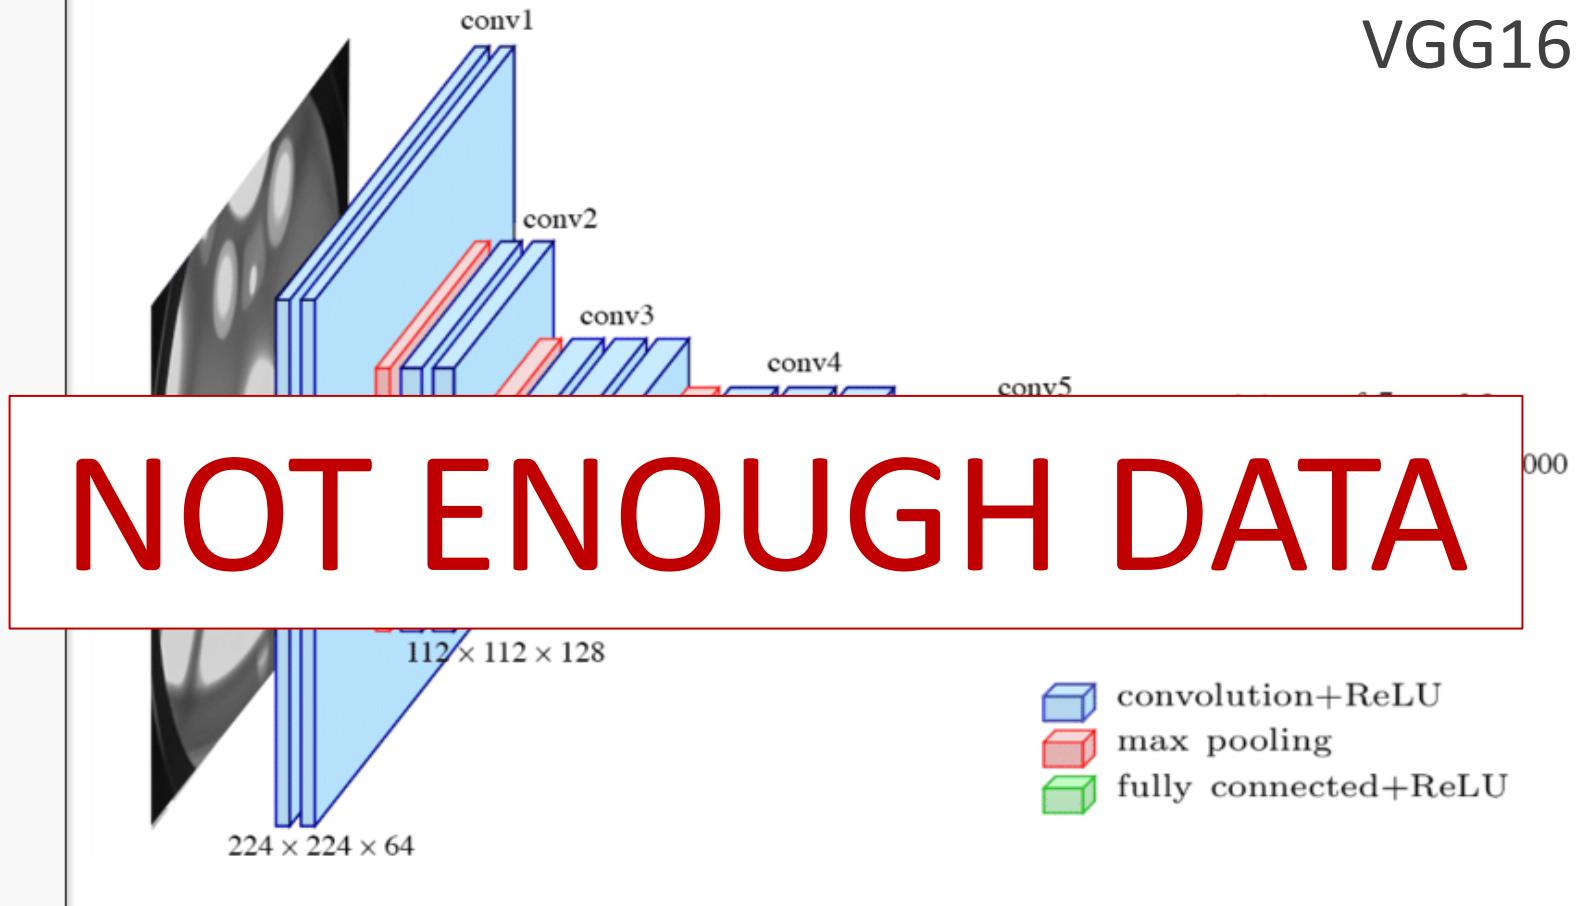
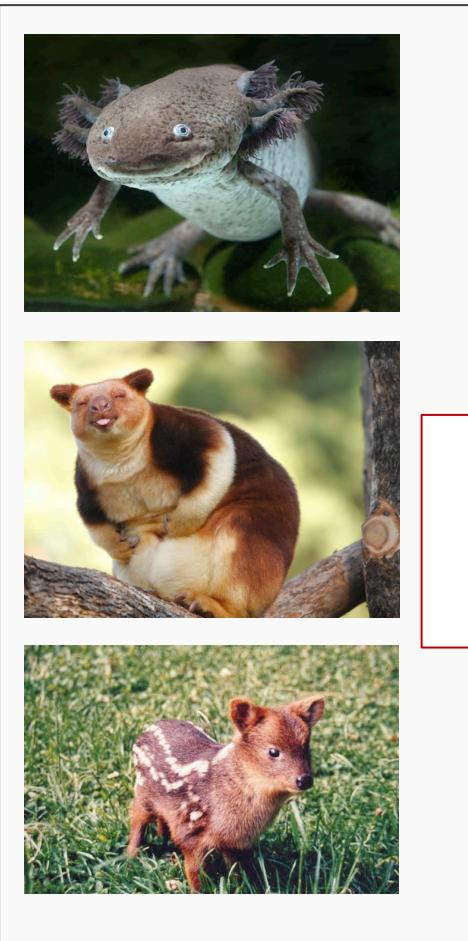


Vincent Casser

# Workshop Overview for Session 2

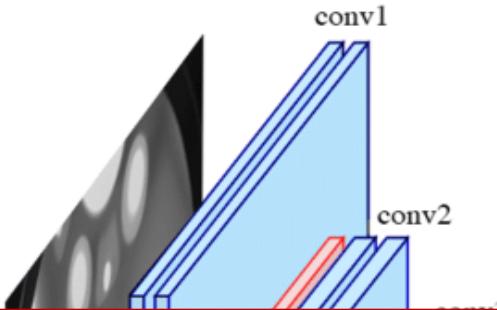
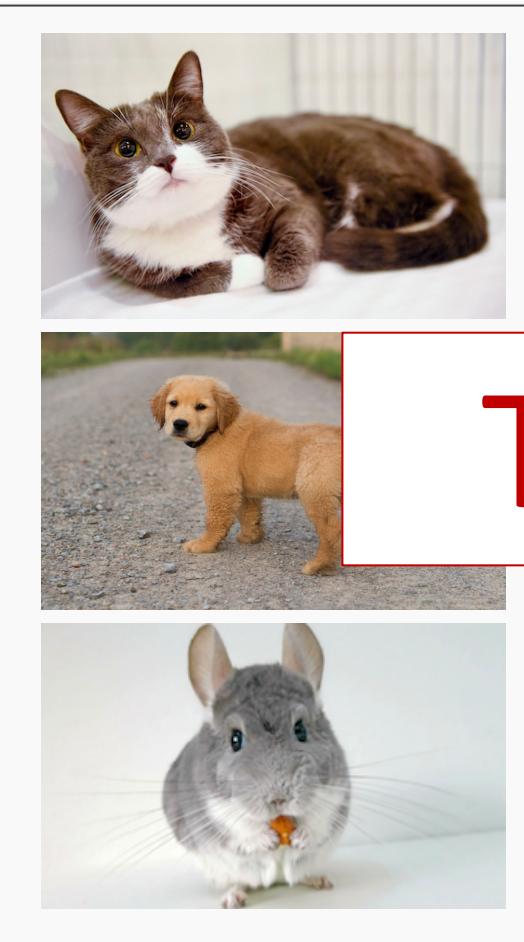


# Classify Rarest Animals



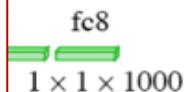
Number of parameters: 134,268,737  
Data Set: Few hundred images

# Classify Cats, Dogs, Chinchillas etc



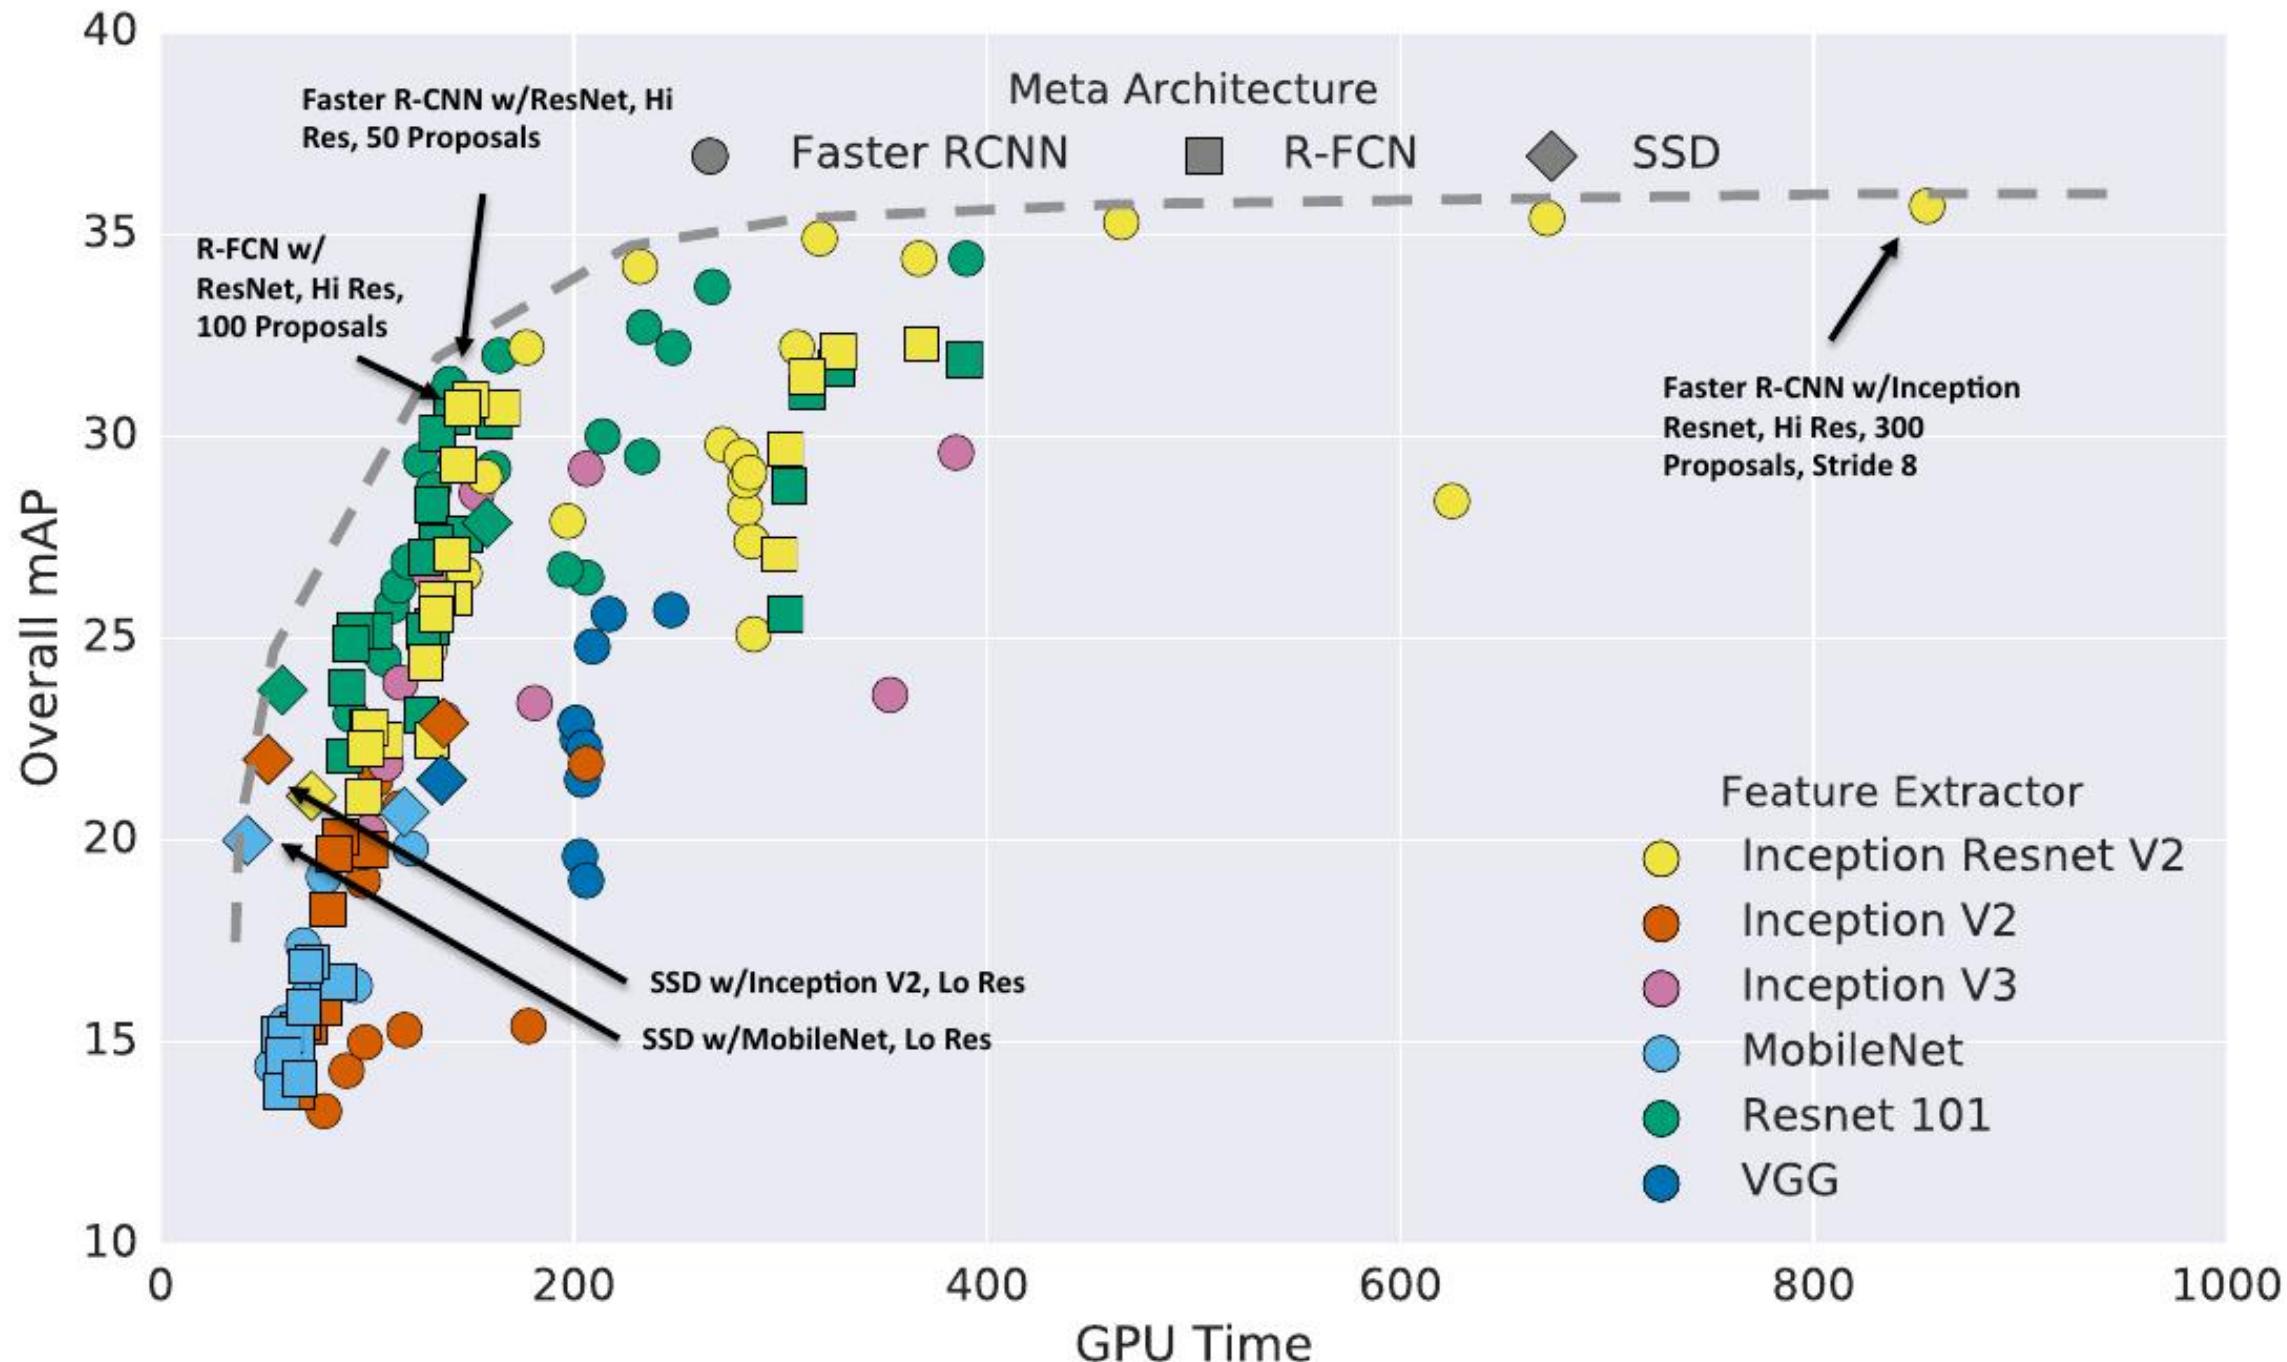
VGG16

**TAKES TOO LONG**



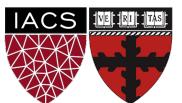
- convolution+ReLU
- max pooling
- fully connected+ReLU

Number of parameters: 134,268,737  
Enough training data. ImageNet approximate 1.2M

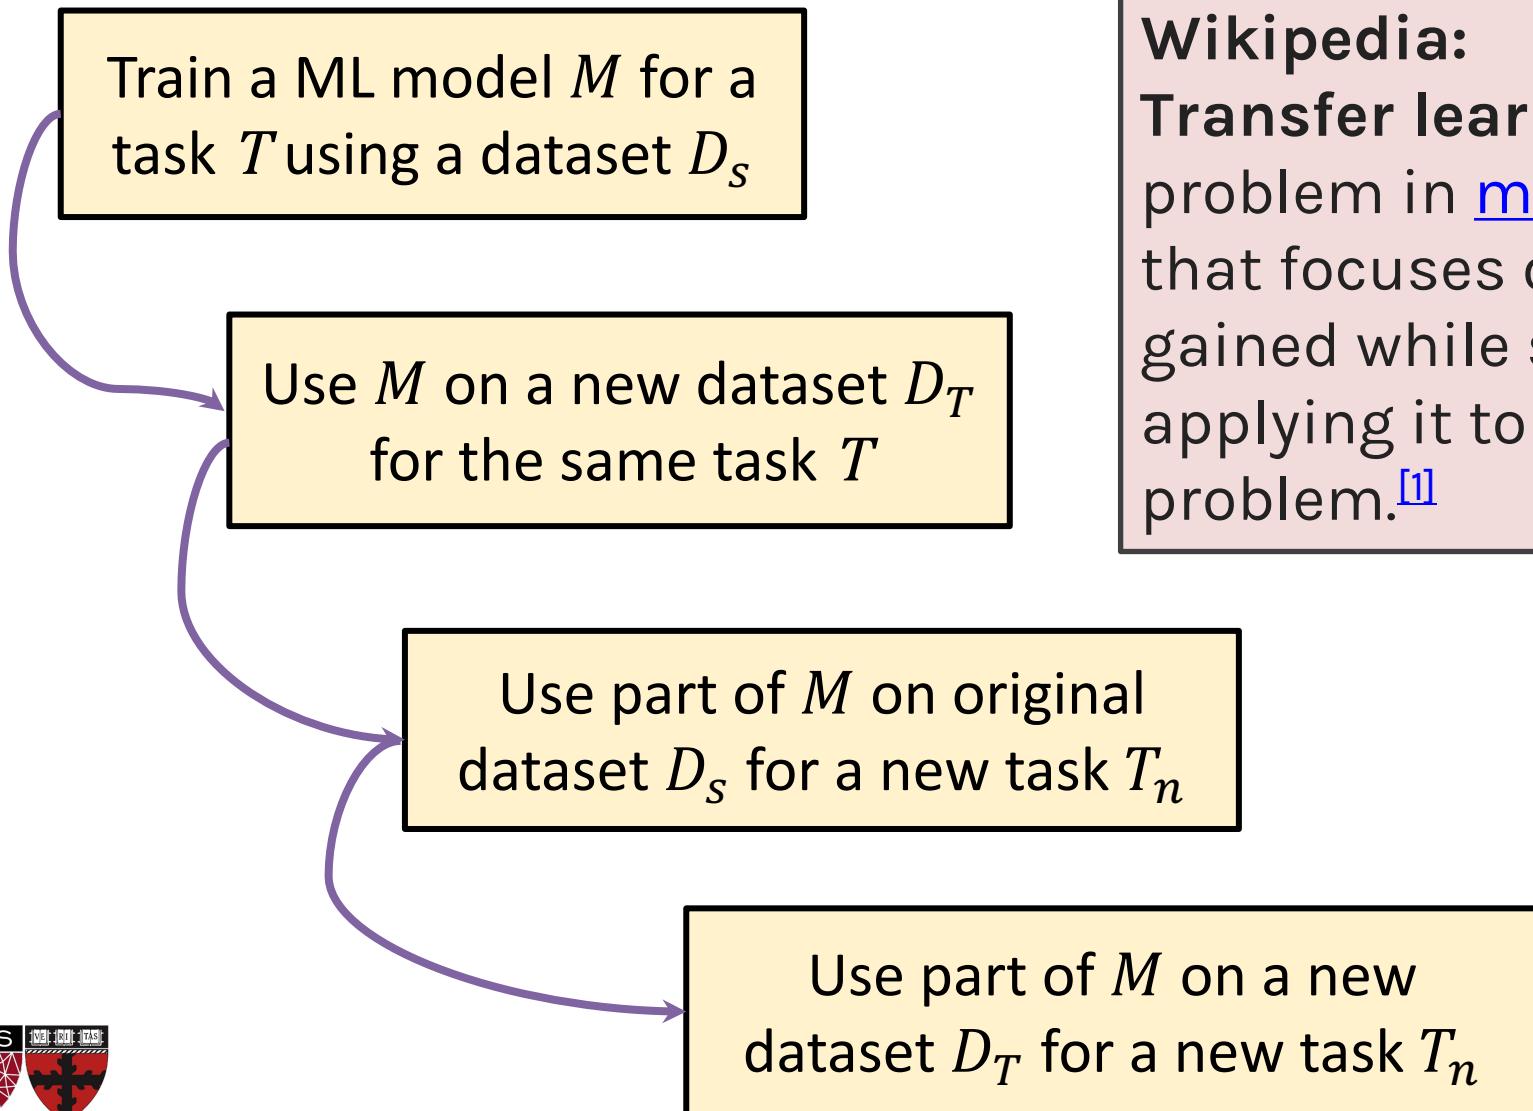


# Transfer Learning To The Rescue

How do you build an image classifier that can be trained in a few minutes on a CPU with very little data?



# Basic idea of Transfer Learning

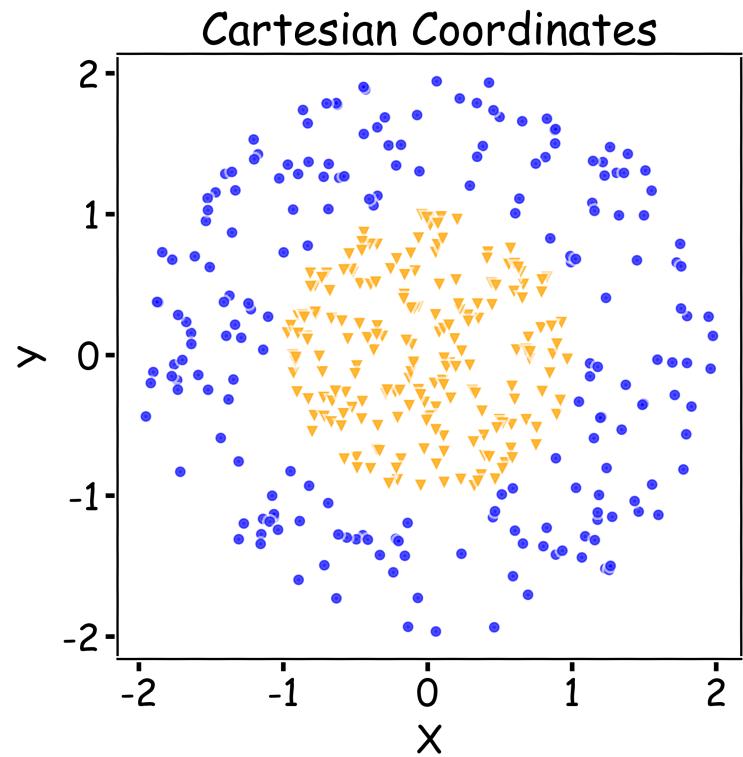


## Wikipedia:

Transfer learning (TL) is a research problem in [machine learning](#) (ML) that focuses on storing knowledge gained while solving one problem and applying it to a different but related problem.<sup>[1]</sup>

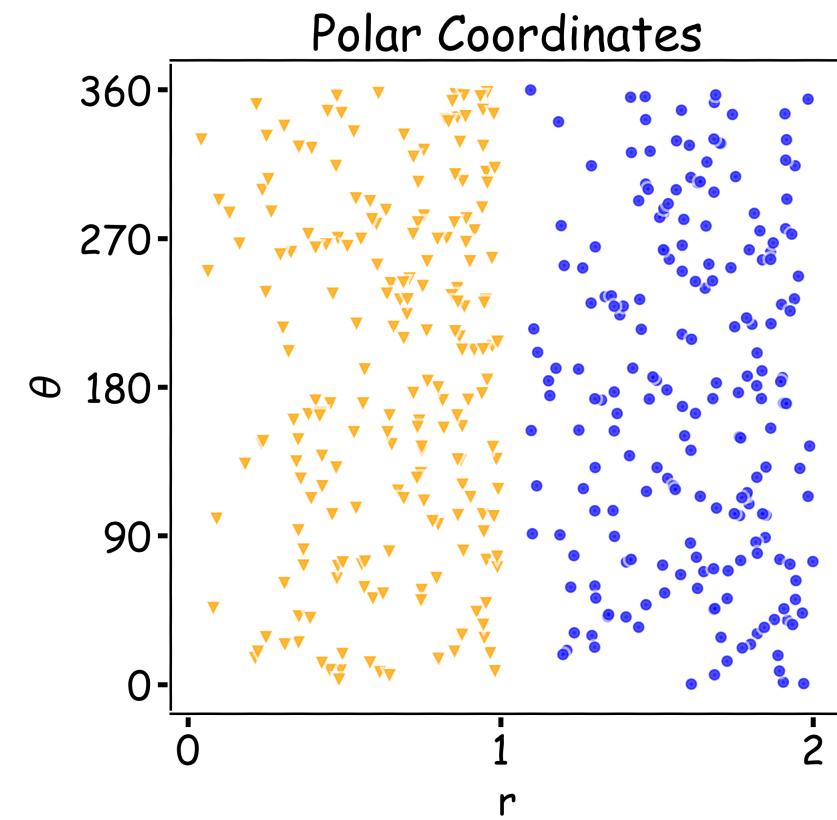
# Key Idea: Representation Learning

Relatively difficult task



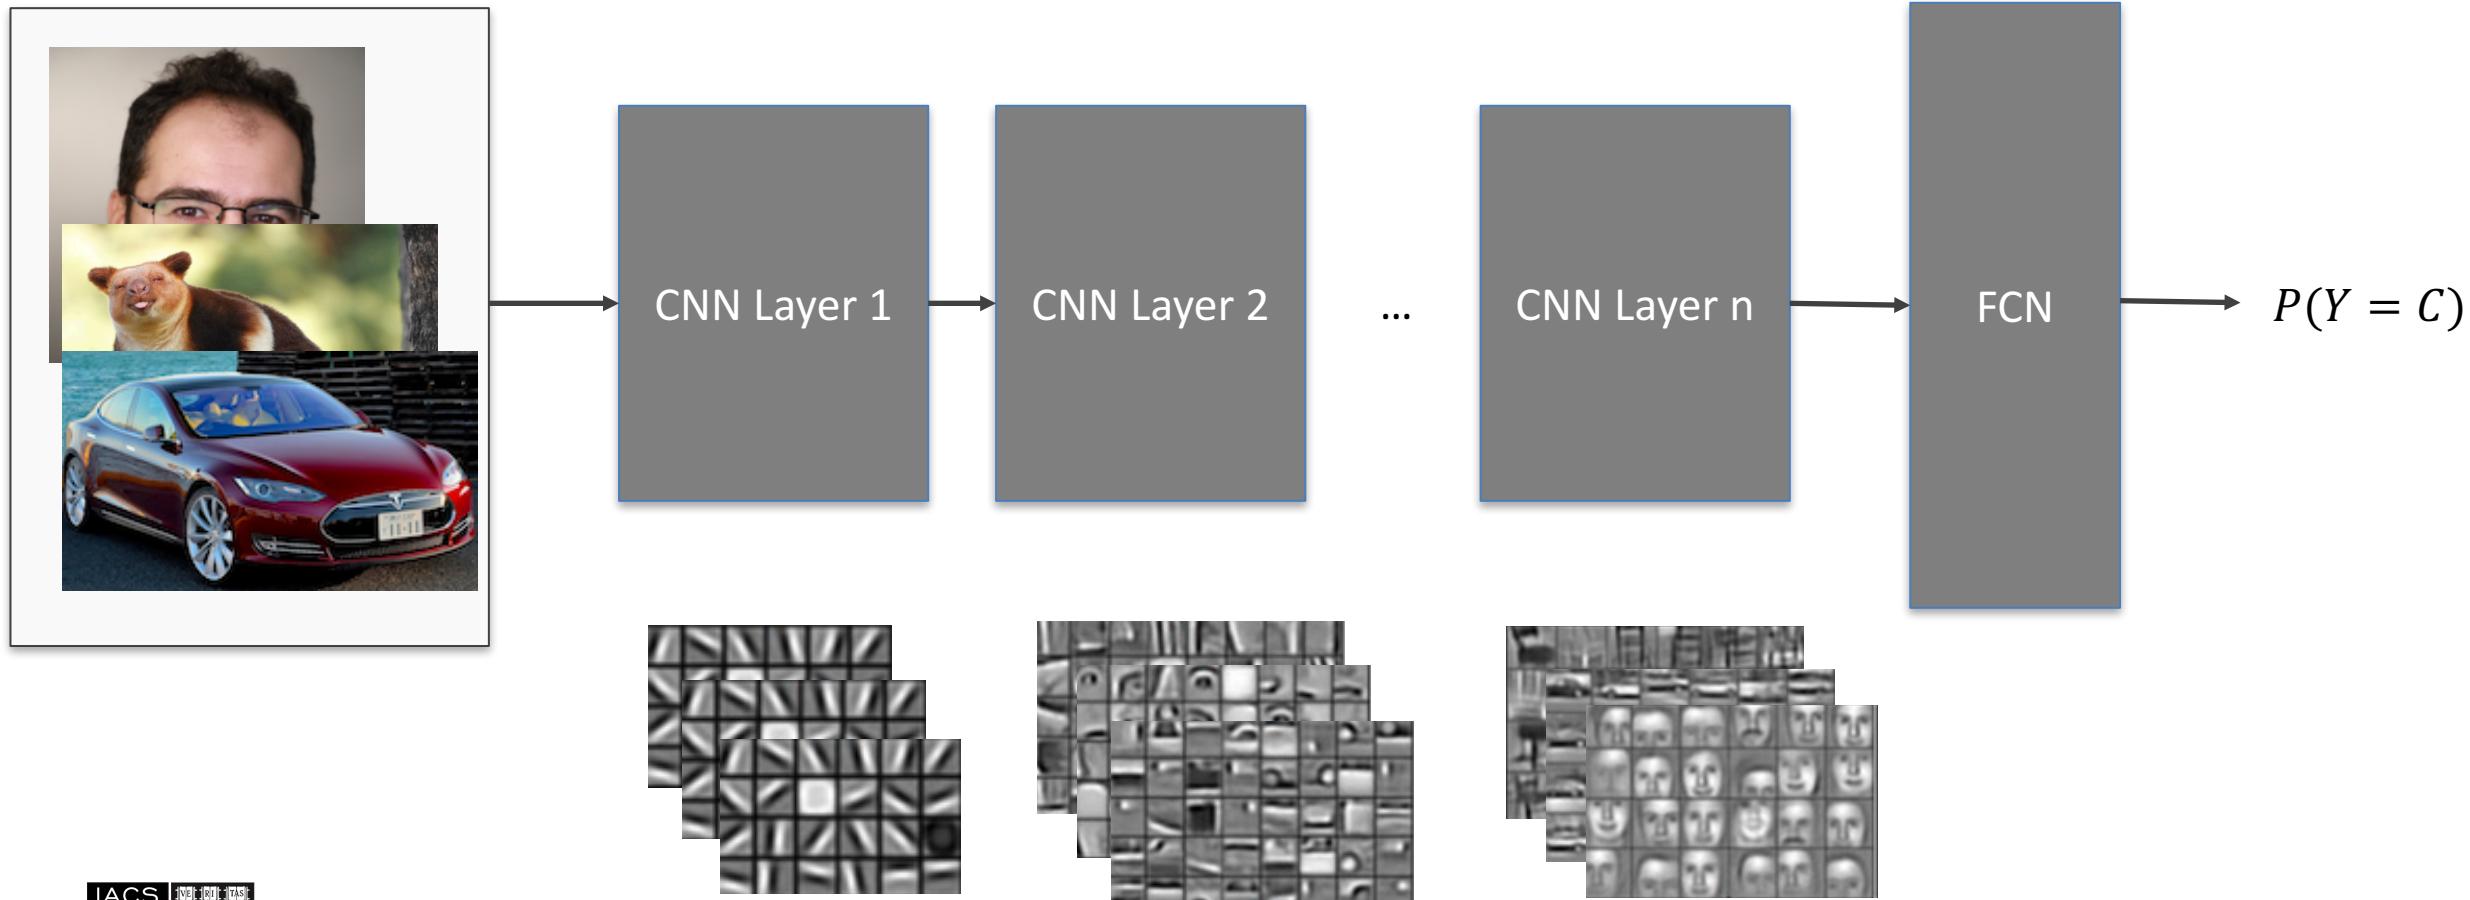
Transform:  
 $(X, Y) \rightarrow (r, \theta)$

Easier task



# Representation Learning

Task: classify cars, people, animals and objects



# Transfer Learning To The Rescue

---

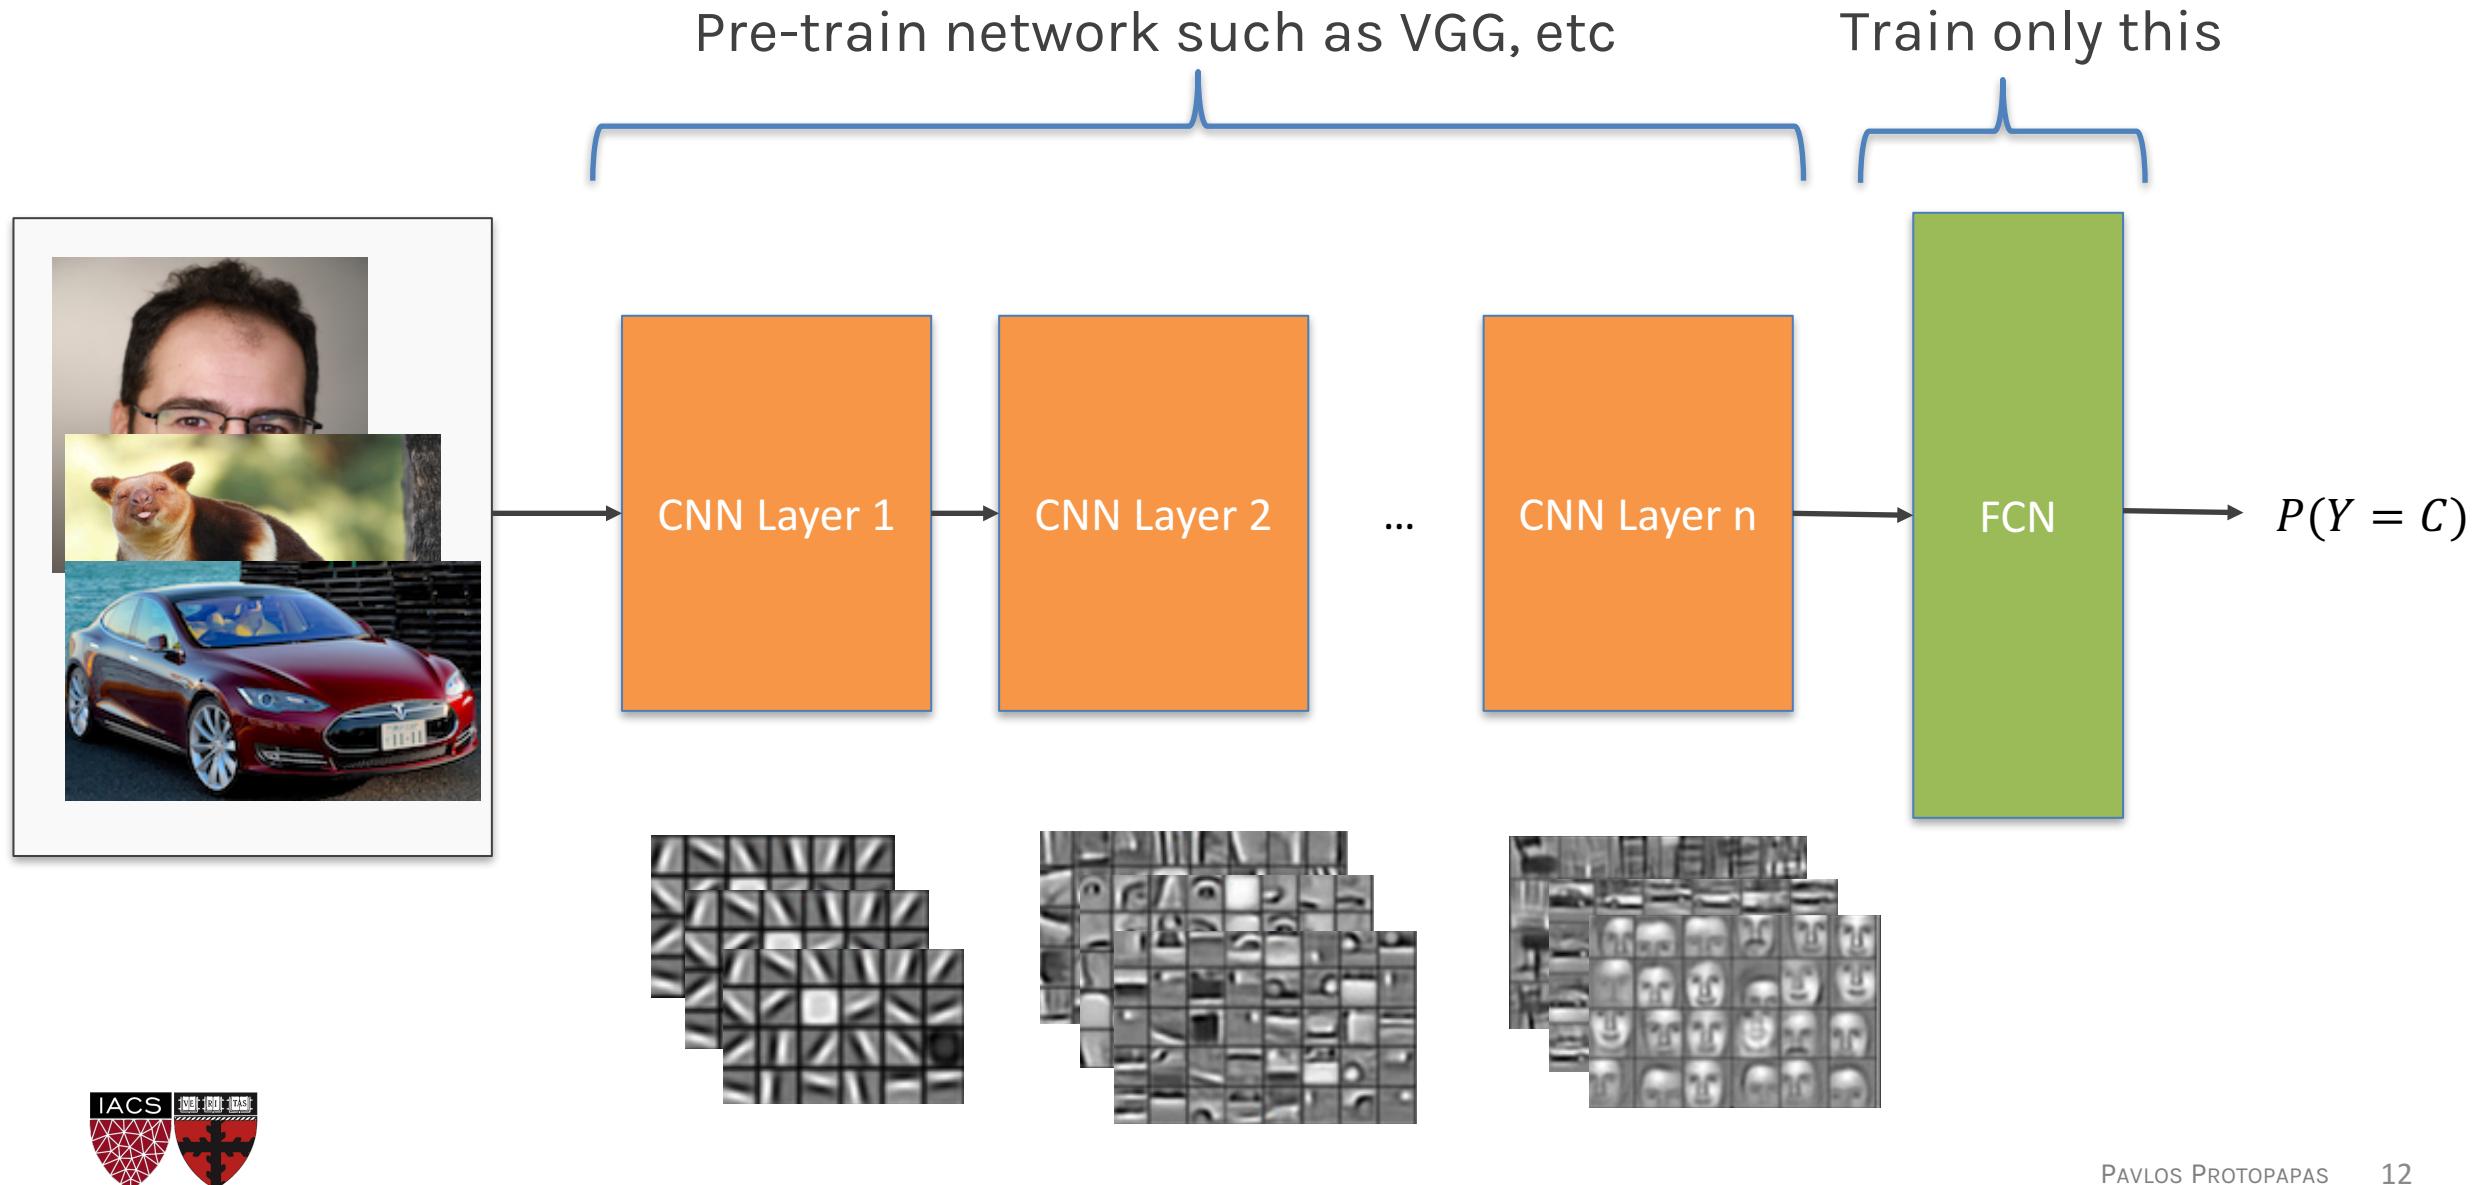
How do you make an image classifier that can be trained in a few minutes on a CPU with very little data?

**Use pre-trained models**, i.e., models with known weights.

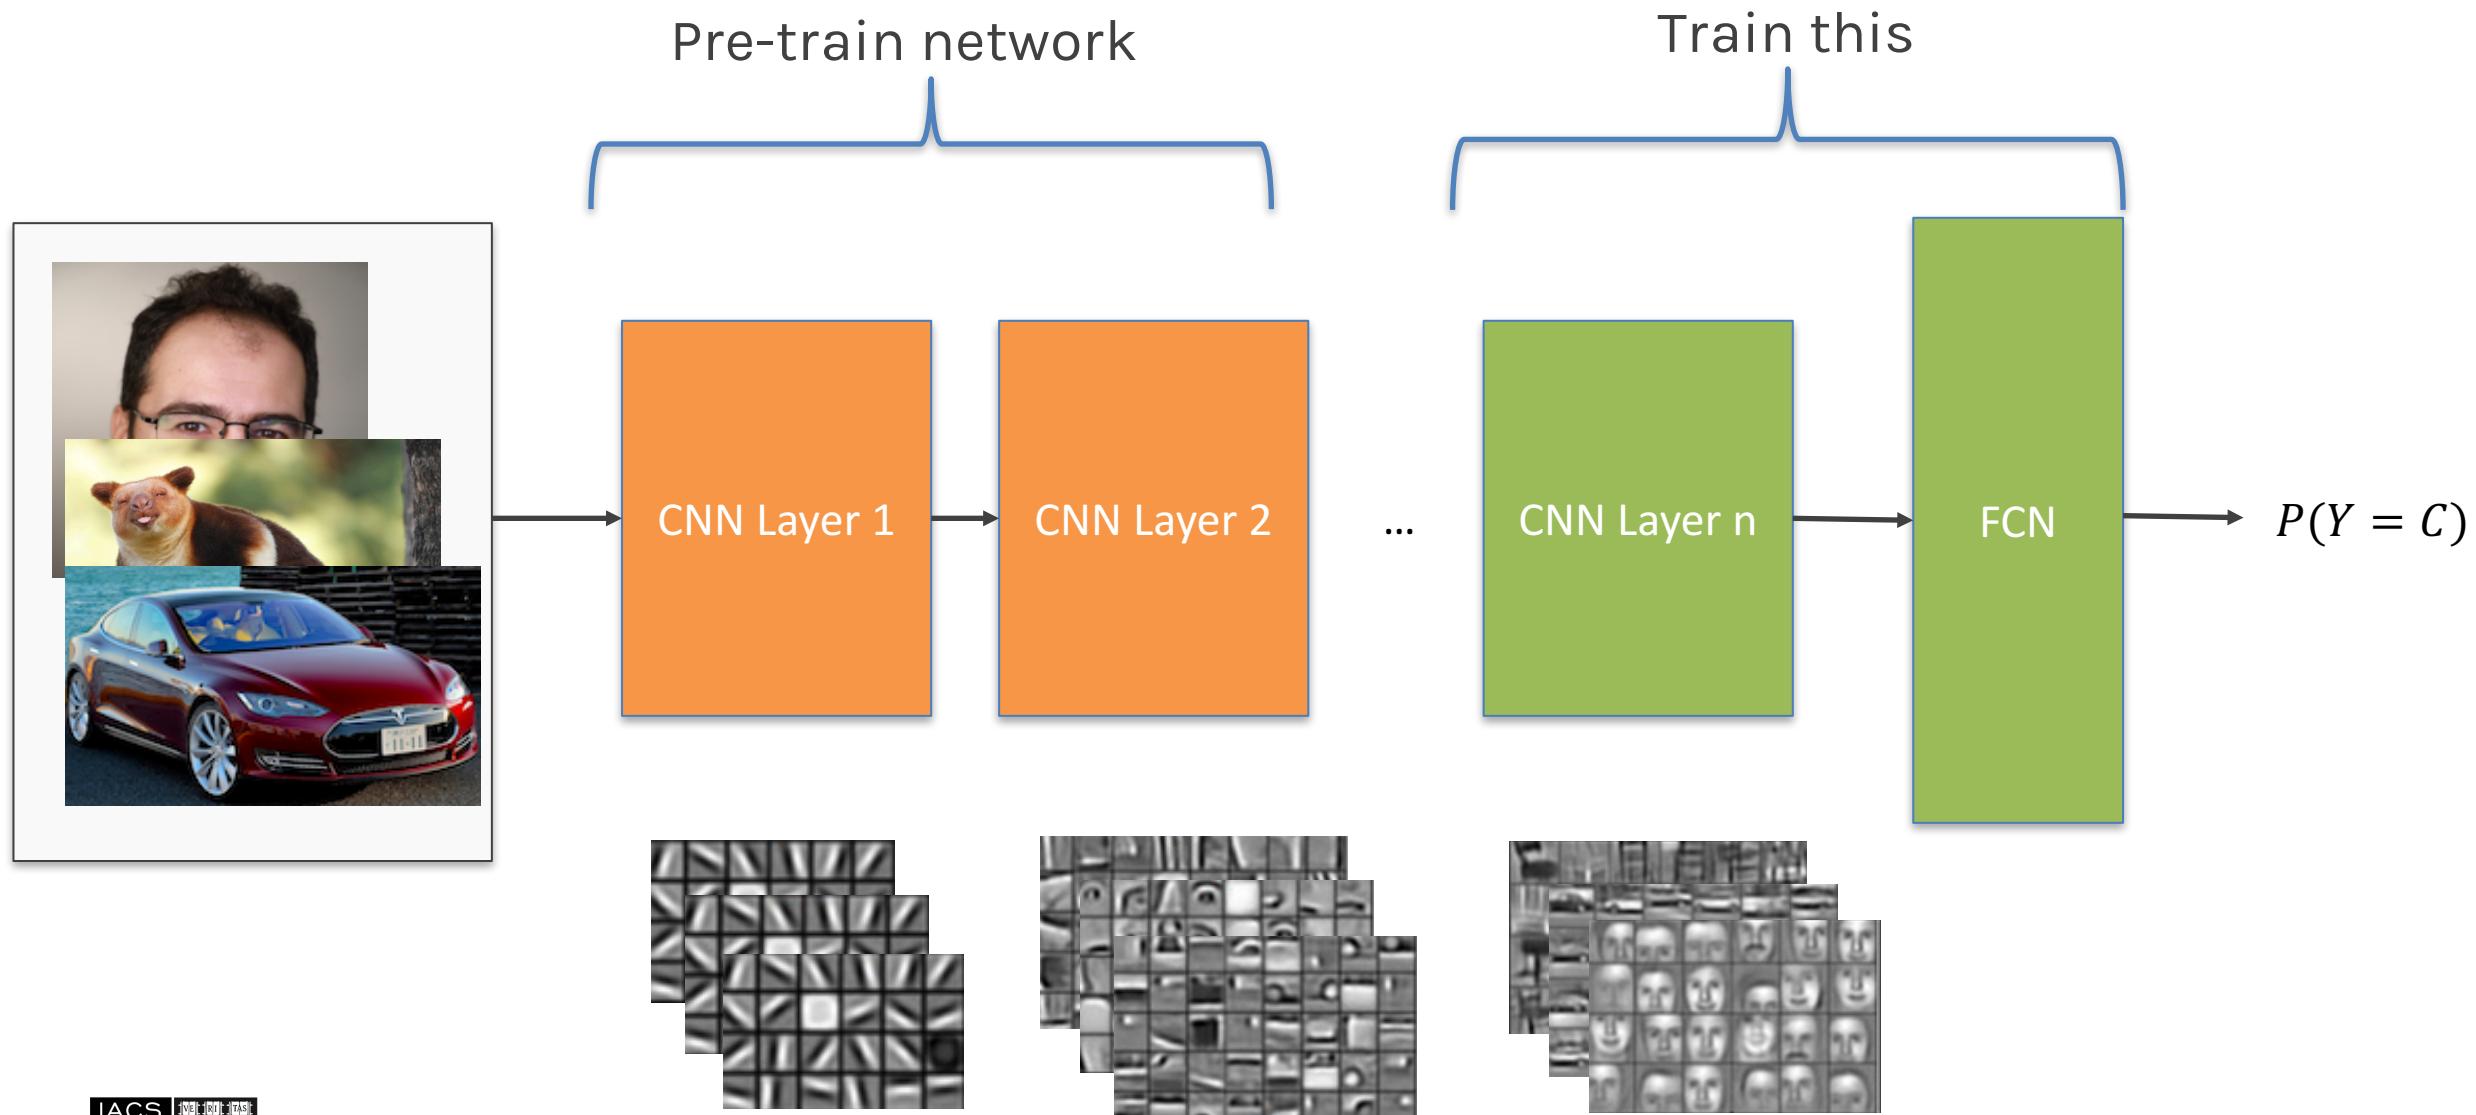
**Main Idea:** earlier layers of a network learn low level features, which can be adapted to new domains by changing weights at later and fully-connected layers.

**Example:** use ImageNet trained with any sophisticated huge network. Then retrain it on a few images

# Transfer Learning

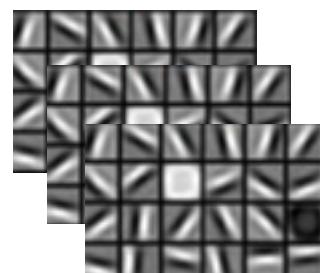
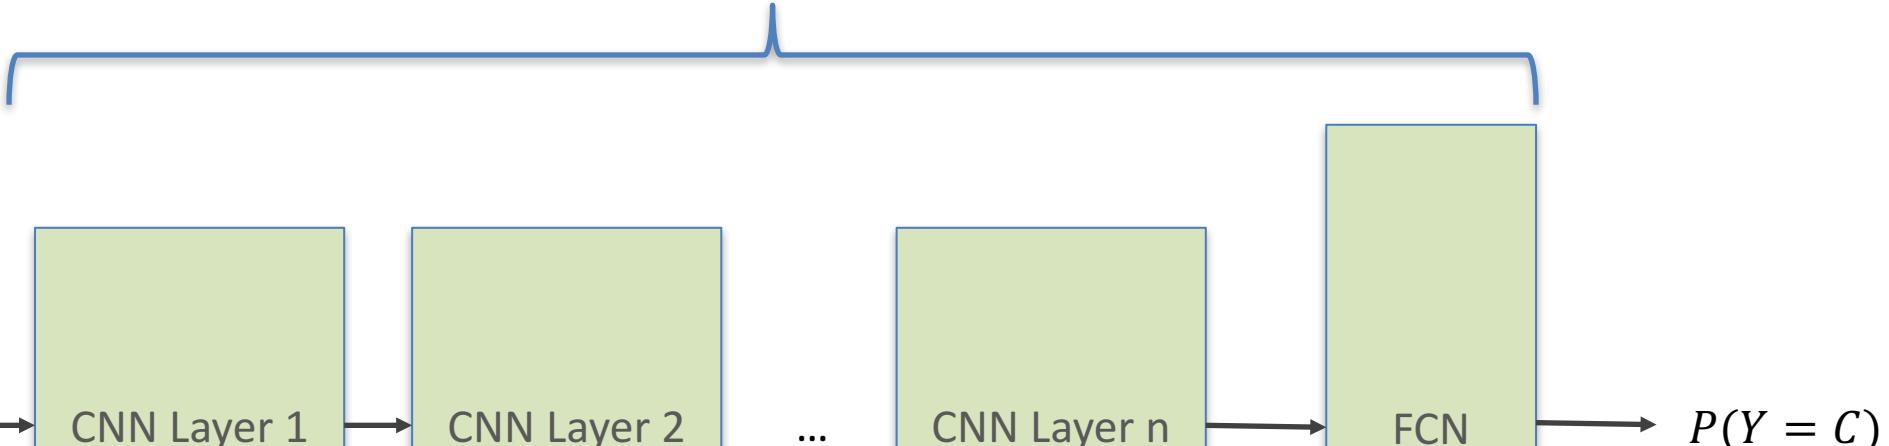
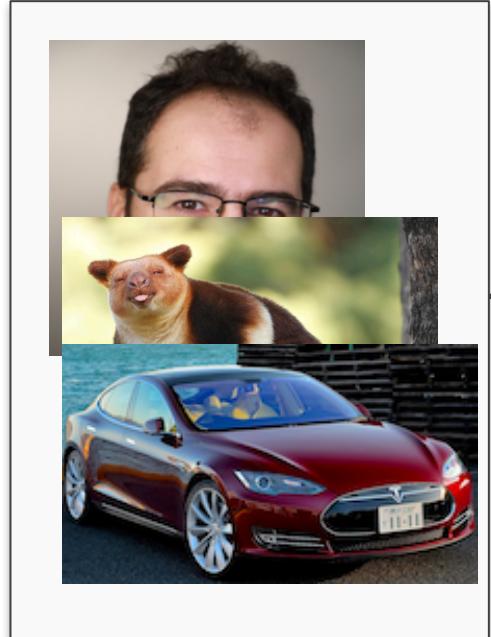


# Transfer Learning



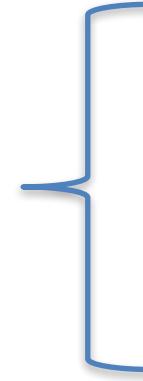
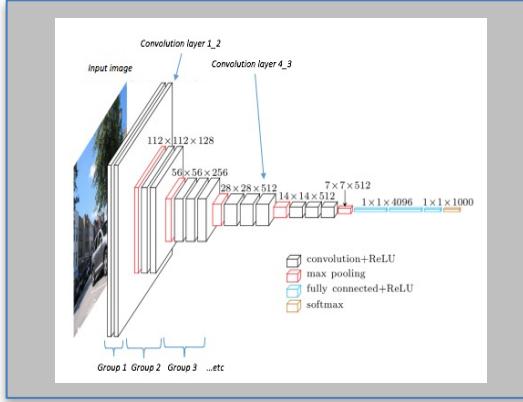
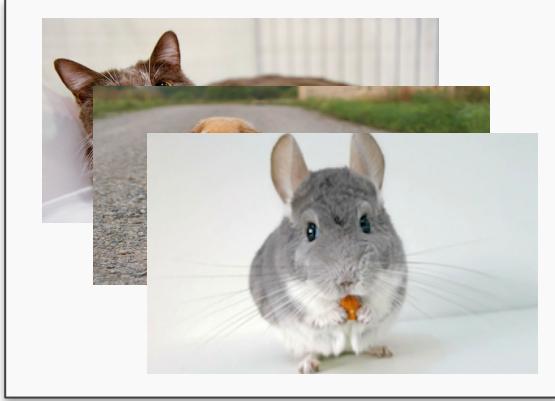
# Transfer Learning - fine tuning

Train everything but start with weights that are trained already



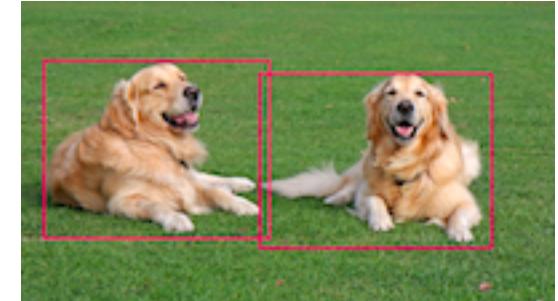
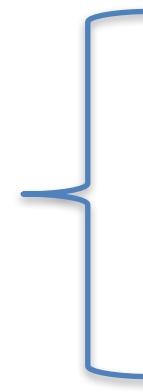
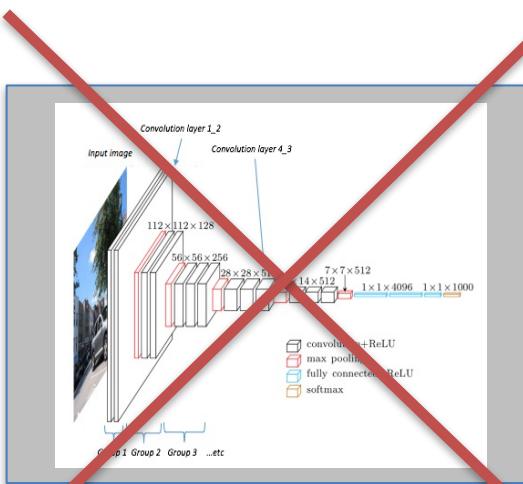
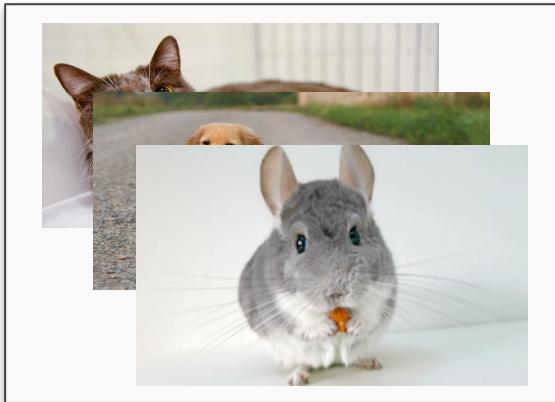
# Transfer Learning Across Tasks

## Classification:



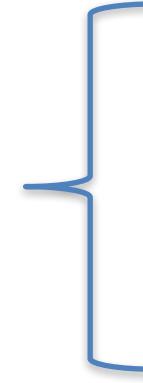
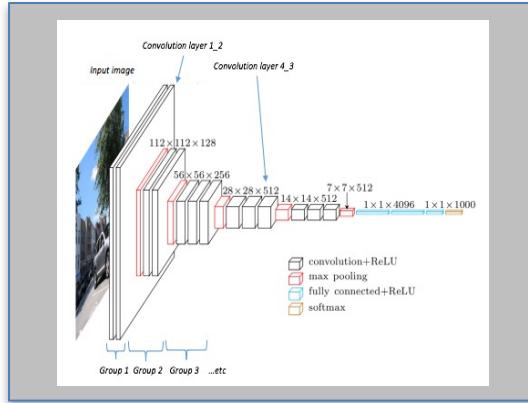
P(cat)  
P(chinchilla)  
P(dog)

## Object Detection:



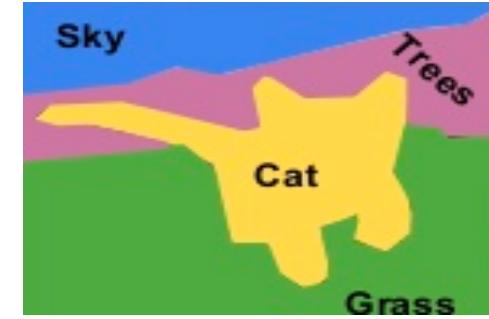
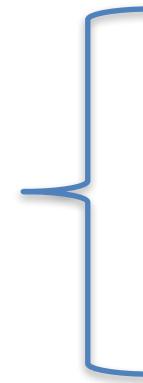
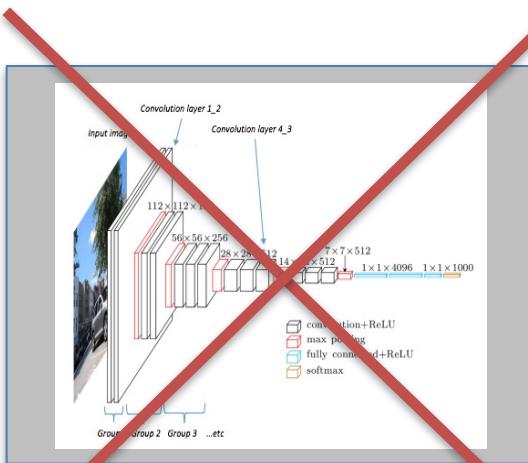
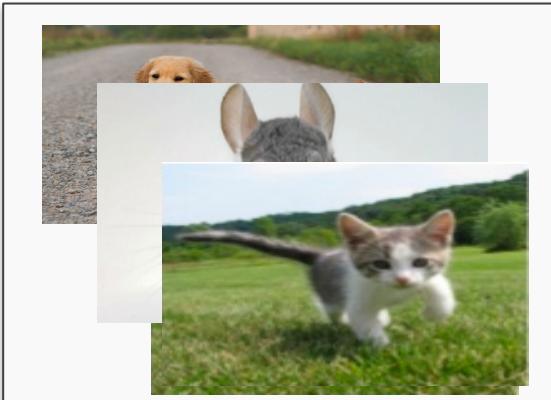
# Transfer Learning Across Tasks (cont)

## Classification:

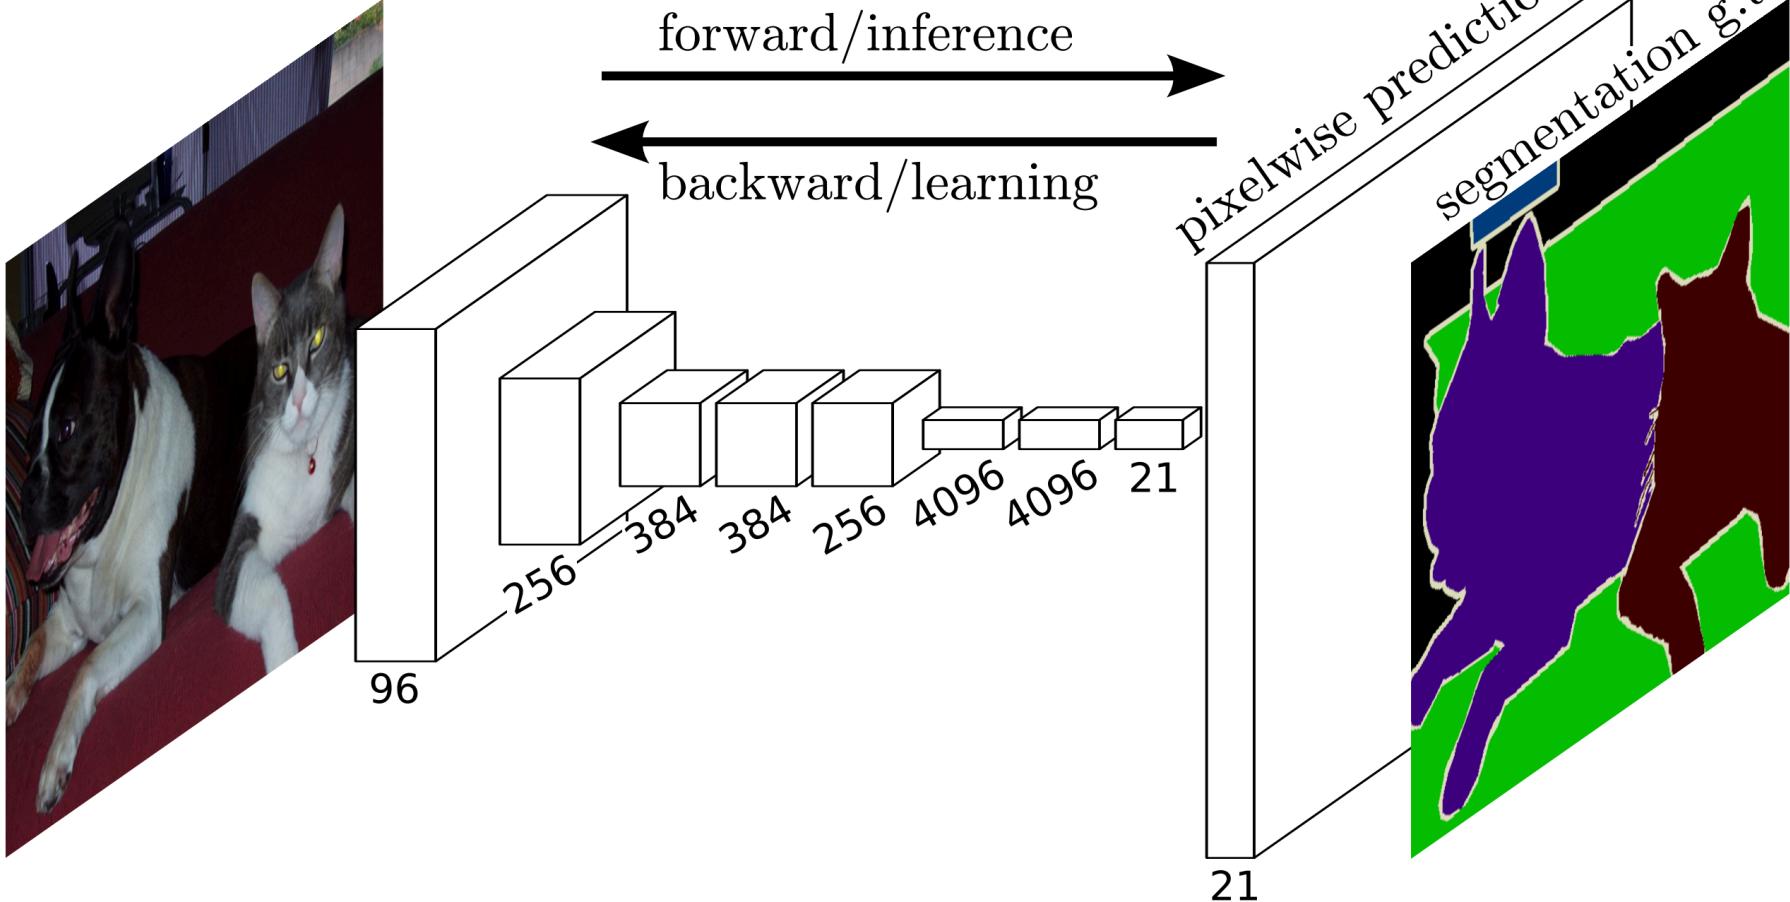


P(cat)  
P(chinchilla)  
P(dog)

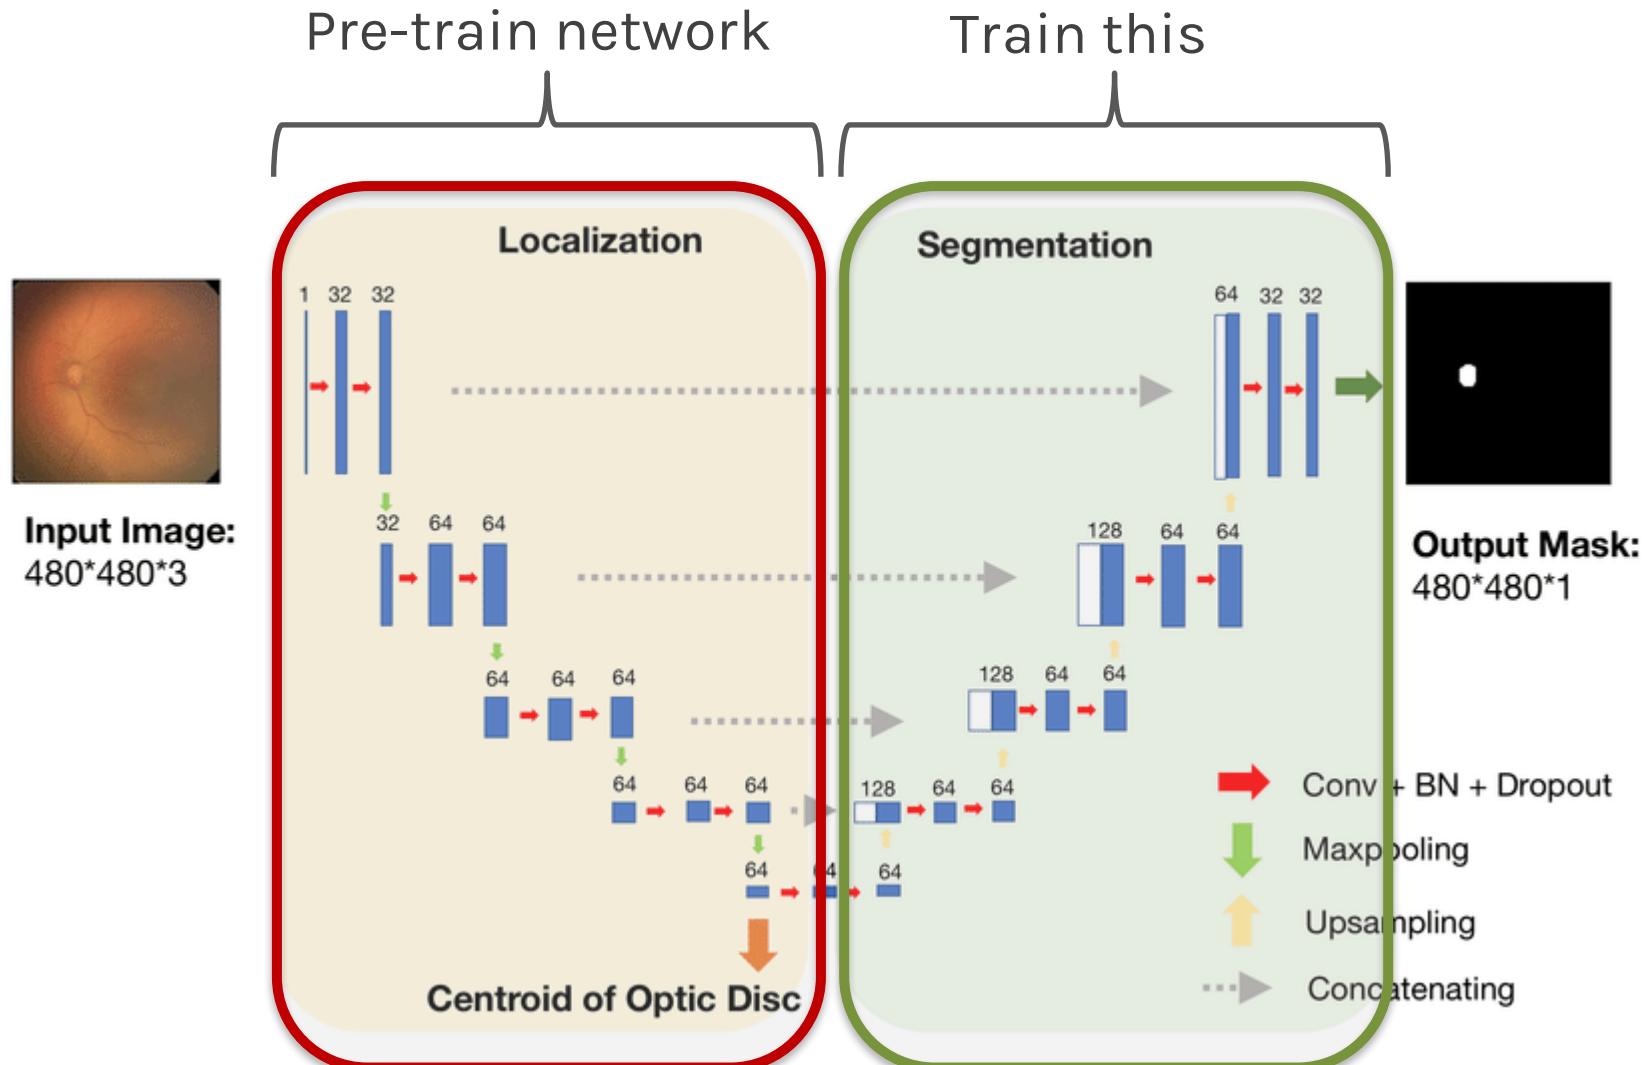
## Segmentation:



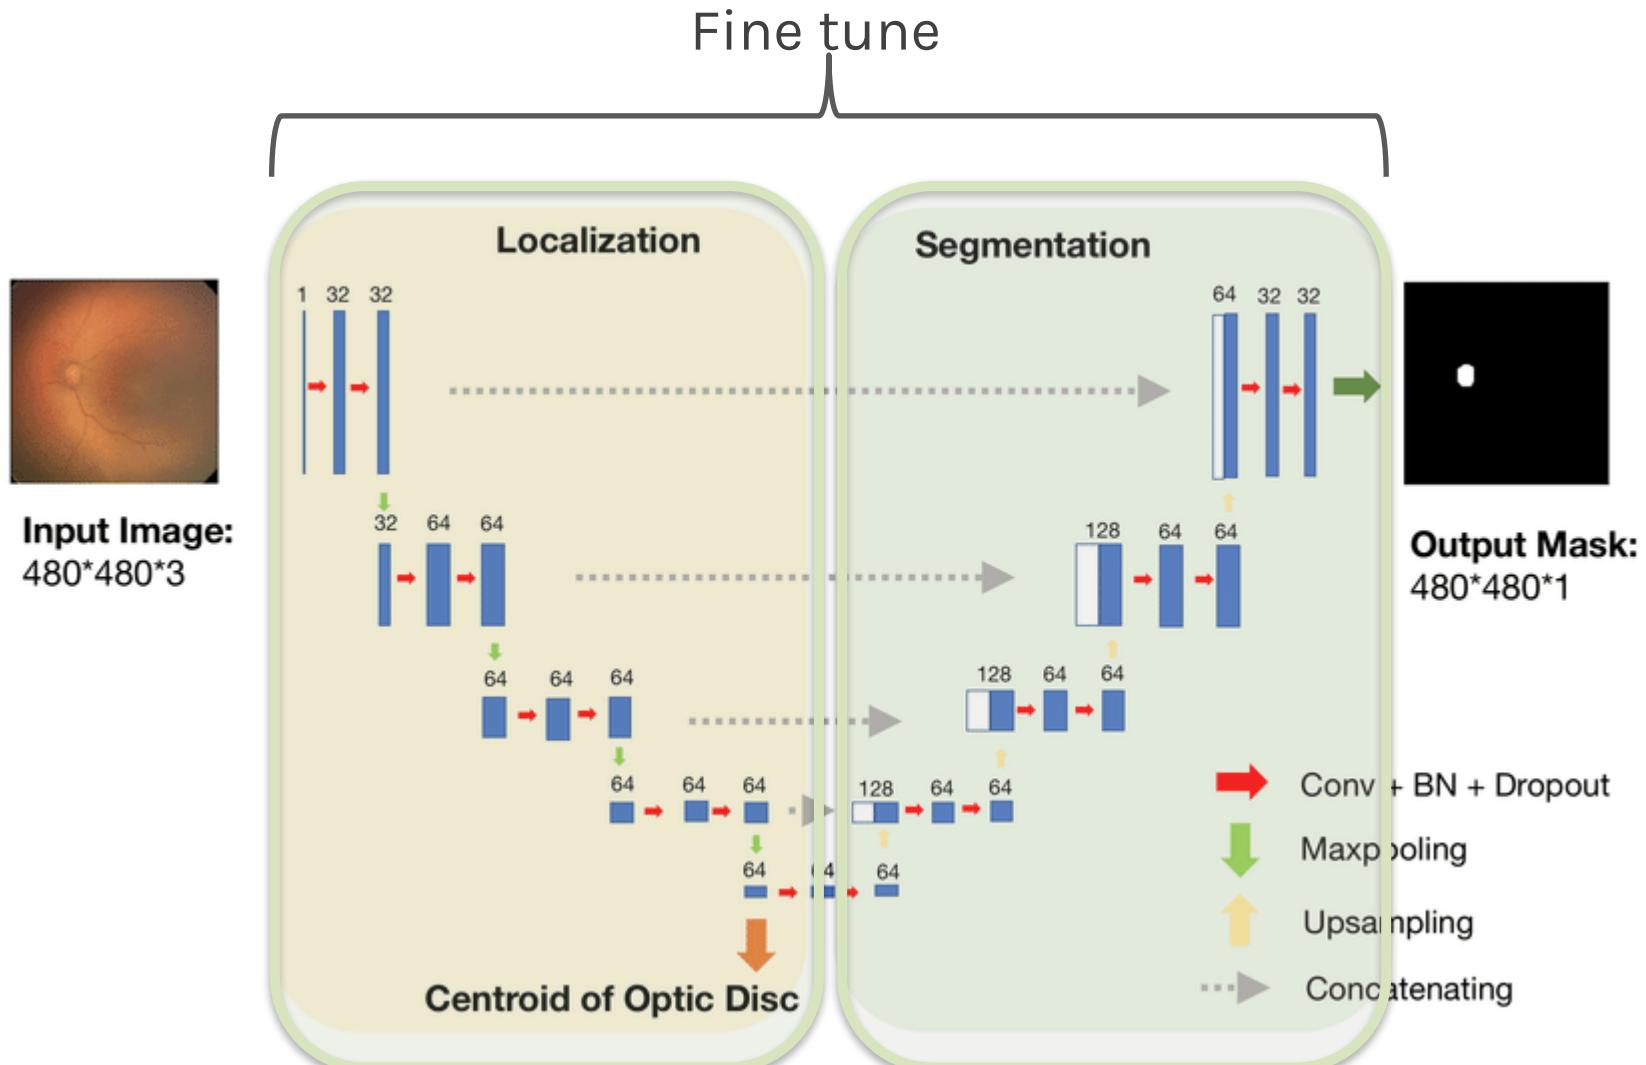
# Segmentation



# Segmentation: U-net architecture

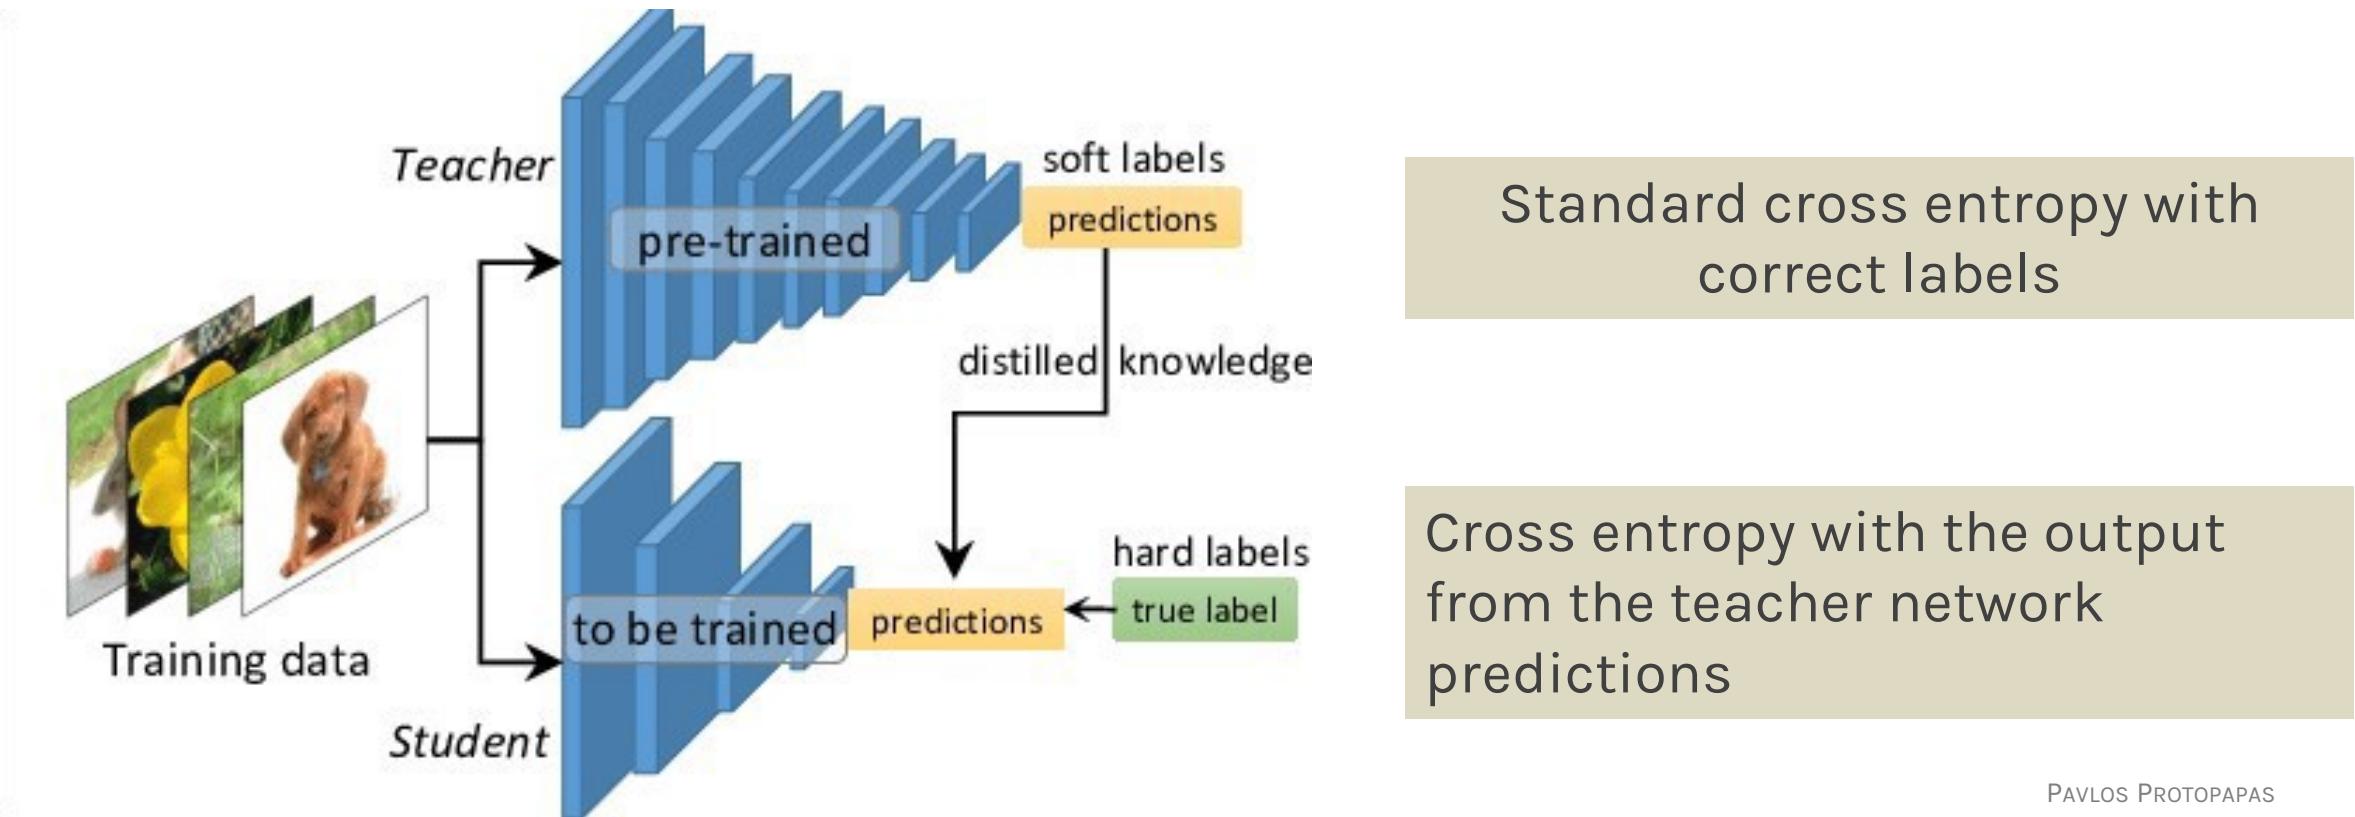


# Segmentation (fine tuning)

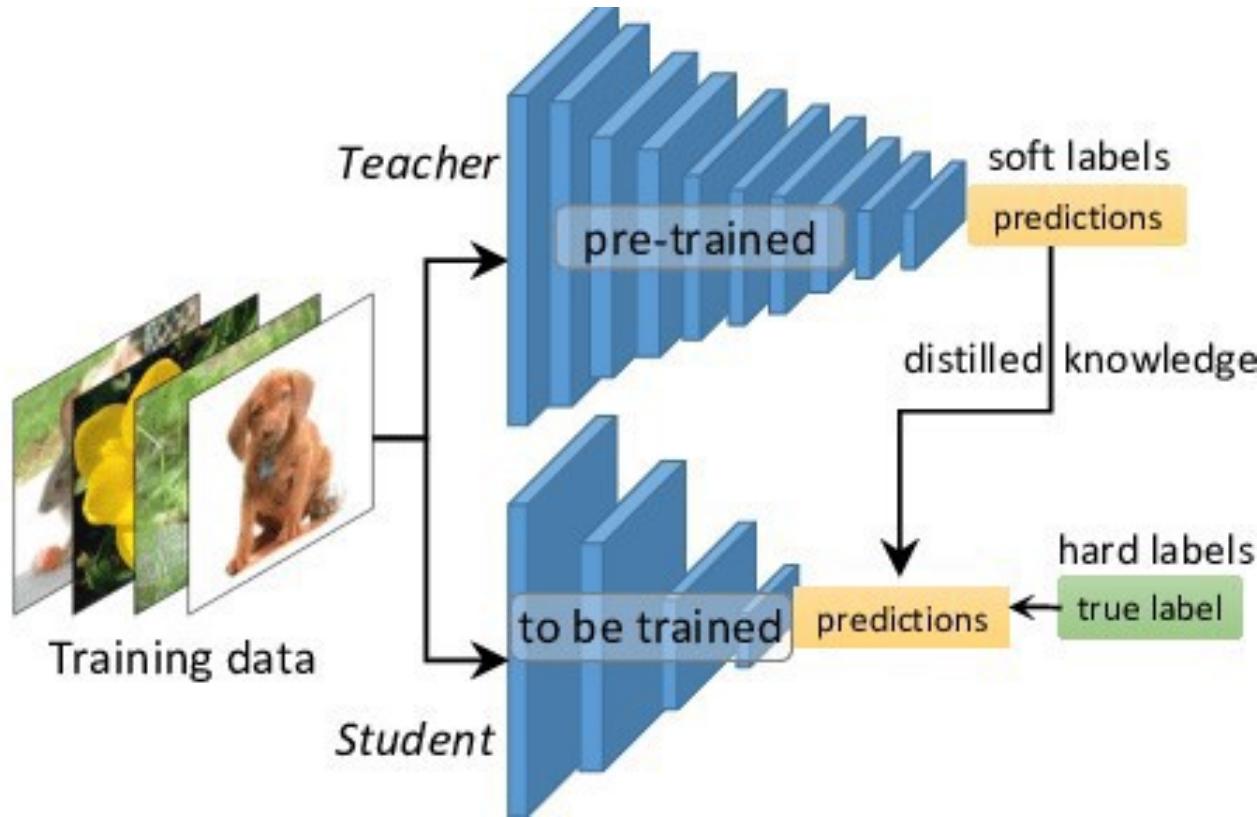


# Transfer Learning through Network Distillation

**Question:** is it possible to **distill** and **compress** the knowledge of the large and complex training model (**the teacher**) into a small and simple deployment model (**the student**)?



# Transfer Learning through Network Distillation (temperature)



Standard cross entropy with  
correct labels

↓  
Soften labels with  
temperature

$$q_i = \frac{e^{z_i/T}}{\sum e^{z_i/T}}$$

Cross entropy with the soft targets  
from the teacher network  
predictions

# Workshop Overview for Session 2

